

What it takes to make RAG work on regulation.

Four AI-governance
documents. 30 questions to
test it, then 45 more it had
never seen.

A model on its own generates from what it learned in training.

It doesn't inherently give you
source-level provenance, and
its knowledge has a cutoff.

On regulation, that gap
matters. An answer you can't
trace back to the source is
hard to trust.

Retrieve first. Then answer only from what came back.

Built over four real-world
governance sources:

EU AI Act	BINDING REGULATION
NIST AI Risk Management Framework	RISK GUIDANCE
NIST Generative AI Profile	GENAI RISK
OWASP LLM Top 10	SECURITY

403

sections parsed from four documents

856

passages, chunked so no rule is split from its conditions

204K

tokens of source material

1.2 MB

the entire search index. No vector database

At this size an exact search across every passage takes 141 milliseconds.

Retrieving more context didn't help the hard questions.

3/8

top 5 passages

3/8

top 10 passages — twice as much, no change

Some answers required following a trail. One section points at another, and neither means anything alone.

Annex III

lists CV screening as high-risk



Article 6

defines the conditions

111 of the 403 sections cite another section like this. In retrieval terms: multi-hop.

Cross-reference questions
answered

7/8

Up from 3/8. The system records every cross-reference when it parses the documents, then follows them at query time.

No agent. No second model call. It can't invent a link that isn't in the text.

24/26

retrieved the evidence needed

4/4

unsupported questions correctly refused

0

answers returned without citations

Refusal is enforced where the answer is written, not by a confidence score. The highest-scoring unanswerable question outscored the lowest-scoring answerable one.

45 QUESTIONS IT HAD NEVER BEEN TESTED AGAINST

One fabricated answer in 45.

Graded by a separate model against the source material, blind to how the system produced the answers.

When retrieval fell short, it declined nine times for every answer it got wrong.

RAG isn't “put documents in a model and ask questions.”

It's the source data. The retrieval logic. The failure modes.

And whether your evaluation is actually testing the system, or quietly validating the choices you already made.