

What it takes to make RAG work on regulation.

Four AI-governance documents, a system that answers only from them, and an evaluation built to locate where it fails.

The model does not answer from what it learned in training.

It answers from passages fetched out of a fixed corpus at query time, and is constrained to those passages.

Which relocates the problem. System quality is no longer model quality — it is **whether retrieval put the right evidence in front of it**, and that is the part this measures.

On regulation, an unciteable answer is not an answer.

The AI Act was amended in July 2026. A model answering from training text cannot tell you which version it learned, or point you at the clause.

Every claim here carries a citation to the article, annex or page it came from, so the reader can verify it rather than trust it.

Four sources that overlap in subject and differ in force.

EU AI Act

BINDING LAW

The regulation. Defines high-risk classification and the obligations that follow from it. Consolidated text, including the 2026 amendments.

NIST AI Risk Management Framework

VOLUNTARY

A structure for risk work — govern, map, measure, manage. Prescribes practice, imposes no obligation.

NIST Generative AI Profile

VOLUNTARY

The same framework narrowed to generative systems, enumerating risks specific to models that produce text and images.

OWASP LLM Top 10

PRACTITIONER

Failure modes observed in deployed LLM applications: prompt injection, sensitive data disclosure, supply chain, excessive agency.

Overlapping subject, different legal force.

Passages from a voluntary framework and a binding regulation are near neighbours in meaning. An answer citing NIST for an obligation the Act imposes is fluent and wrong.

Which is why the citation stays attached to the claim rather than being summarised away: the reader can see which instrument the answer rests on.

Built once. Then run on every question.

BUILT ONCE, OFFLINE

- 01 **Parse each document into citable sections**
Articles, annexes, numbered sections — the unit a reader would cite

- 02 **Split into passages**
Boundaries chosen so a rule is never severed from its conditions

- 03 **Embed each passage**
Local model, CPU. The index is a file that ships with the application

PER QUESTION

- 04 **Rank every passage by semantic similarity**

- 05 **Expand: the article's opening, then its cross-references**
The two deterministic steps that carry the work

- 06 **Answer with citations, or abstain**

Some provisions are split across two sections, and neither is sufficient alone.

Annex III

lists CV screening as high-risk

**Article 6**

states the conditions that apply

Retrieving more passages does not surface the second. It is not semantically near the question — it is **referenced by the first**.

Cross-references are extracted at parse time and followed at query time.

When a retrieved passage cites another section, that section is pulled into context with it.

Deterministic: no agent, no planner, no second model call, and no capacity to invent a link that is absent from the text.

Similarity search always returns its nearest passages.

However far away they are. There is no native signal for “this corpus does not contain the answer.”

Measured here, retrieval confidence could not separate answerable questions from unanswerable ones. Abstention is therefore enforced at generation, against the retrieved text, rather than by thresholding a score.

Any RAG system will produce a fluent answer to any question.

Fluency is not evidence that retrieval succeeded. A demonstration answer establishes nothing about the system that produced it.

The questions that matter: did it retrieve the sections the answer requires, did it ground each claim in them, and did it abstain when it could not.

HELD-OUT SET · BLIND GRADING

Retrieval parameters were selected against the same set used to report results.

Every figure was in-sample. So: a second set, written from sections the first never touched, keys committed before execution, and both sets graded by a separate model blind to the pipeline.

Scores fell. The independent grader also surfaced a defect in my own refusal detection that self-review had not.

RAG isn't “put documents in a model and ask questions.”

It's the source data. The retrieval logic. The failure modes.

And whether the evaluation is testing the system, or **confirming the choices already made**.

Method, figures, and the failures that remain — including one question it has never retrieved — are in the repository.